

UMassAmherst

COMPSCI 383

Artificial Intelligence

February 5, 2020

Lecture 4

Adversarial Search



Announcements & Administtrivia

- Homework 1 due 11:59pm on Thursday, 2/13.
- Homework 0 is grades will be released soon; please submit a regrade request if you feel you have been wronged.
- Check the Moodle page for updated times/locations of UCA/TA office hours.
- Today I have to leave immediately after class.

Our Time Today

- Understanding how machines can “play” games
- Game playing as search in adversarial environments
- Key concepts: game trees, Min and Max players, utilities of game states
- Methods for searching in games: minimax algorithm, alpha-beta pruning, approximate evaluation functions
- Extending to games with chance elements



Some Context

- In previous lectures, we've talked about searching state spaces for a goal state, by considering available state transitions.
- In many cases, the state space comprises the possible configurations of some system or object.
- We use search to select the best transitions for our agent to reach a goal state from some initial state.
- The search procedure produces a solution path made up of states and legal transitions.

Searching in Games

- For games, the **state** is the total configuration of a game in progress. The successor function defines the legal moves given the game state.
- For the types of games we're going to talk about, there is more than one agent selecting state transitions.
- There may be multiple, equivalent goal states. Furthermore, our goal states are not the same as our opponent's!
- We will identify and order goal states using a **utility function**, which gives us the “value” of any given state toward winning the game.

GREETINGS PROFESSOR FALKEN.
SHALL WE PLAY A GAME?



Why are games interesting to AI? ⁷

- Simple to represent and reason about
- Must consider the moves of an adversary
- Russell & Norvig say:
“Games, like the real world, therefore require the ability to make *some* decision even when calculating the *optimal* decision is infeasible.”

We're Bad at Games

- Chinook was declared the Man-Machine World Champion in checkers (1994)
- Deep Blue defeated the reigning world chess champion Garry Kasparov (1997)
- Google's DeepMind AlphaGo defeated the world's number one Go player Ke Jie (2017)





Types of Games


	deterministic	stochastic
perfect info	chess, checkers, tic-tac-toe, go, othello	backgammon, monopoly
partial info	battleship	poker, scrabble, bridge

Basic Games

- 2-players, alternating moves
- Zero-sum: my gain is your loss!
(Cooperative, non-zero sum games apply game theory)
- Perfect information: both players know the complete game state
(Hidden information games like Poker require belief states)
- Deterministic: no elements of chance (for the moment)

Example: Nim

- On your turn, pick up 1 or 2 coins
- Whoever picks up the last coin loses



Click time!

You are Player 2 in a game of 5-Nim. Before your opponent starts the game, you do some thinking about strategy. Which of the following statements is true?



- (A) You can only win if Player 1 does not act rationally.
- (B) You cannot win if Player 1 selects one coin.
- (C) You can only win if Player 1 selects two coins.
- (D) Your chances of winning are the same no matter how many coins Player 1 selects.

AI Game Terminology

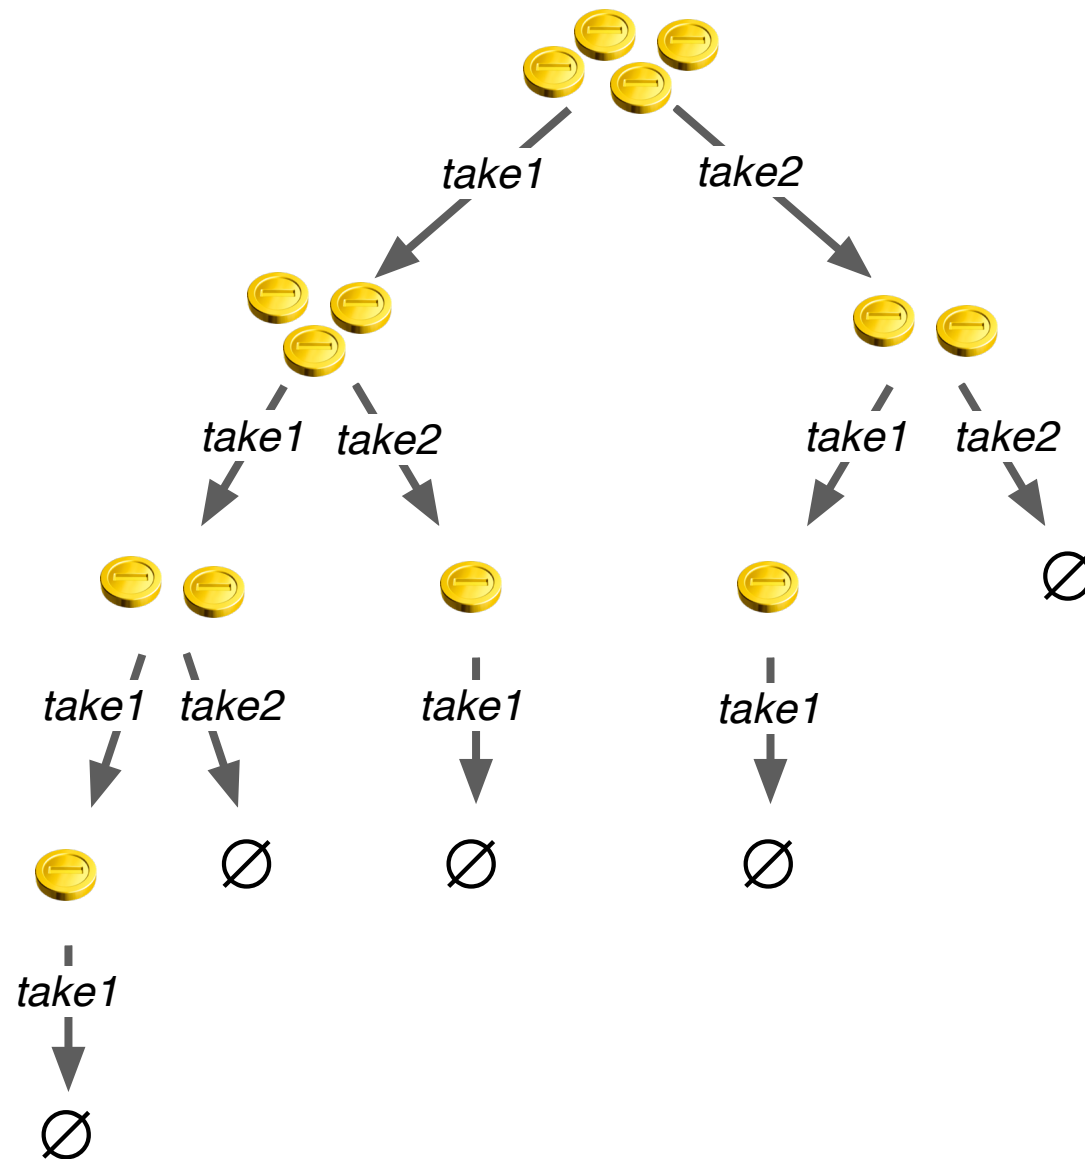
Adversarial Search the process of finding optimal state transitions in the presence of other agents who are working against us

Game Tree data structure defined by the initial game state and the legal moves for each player

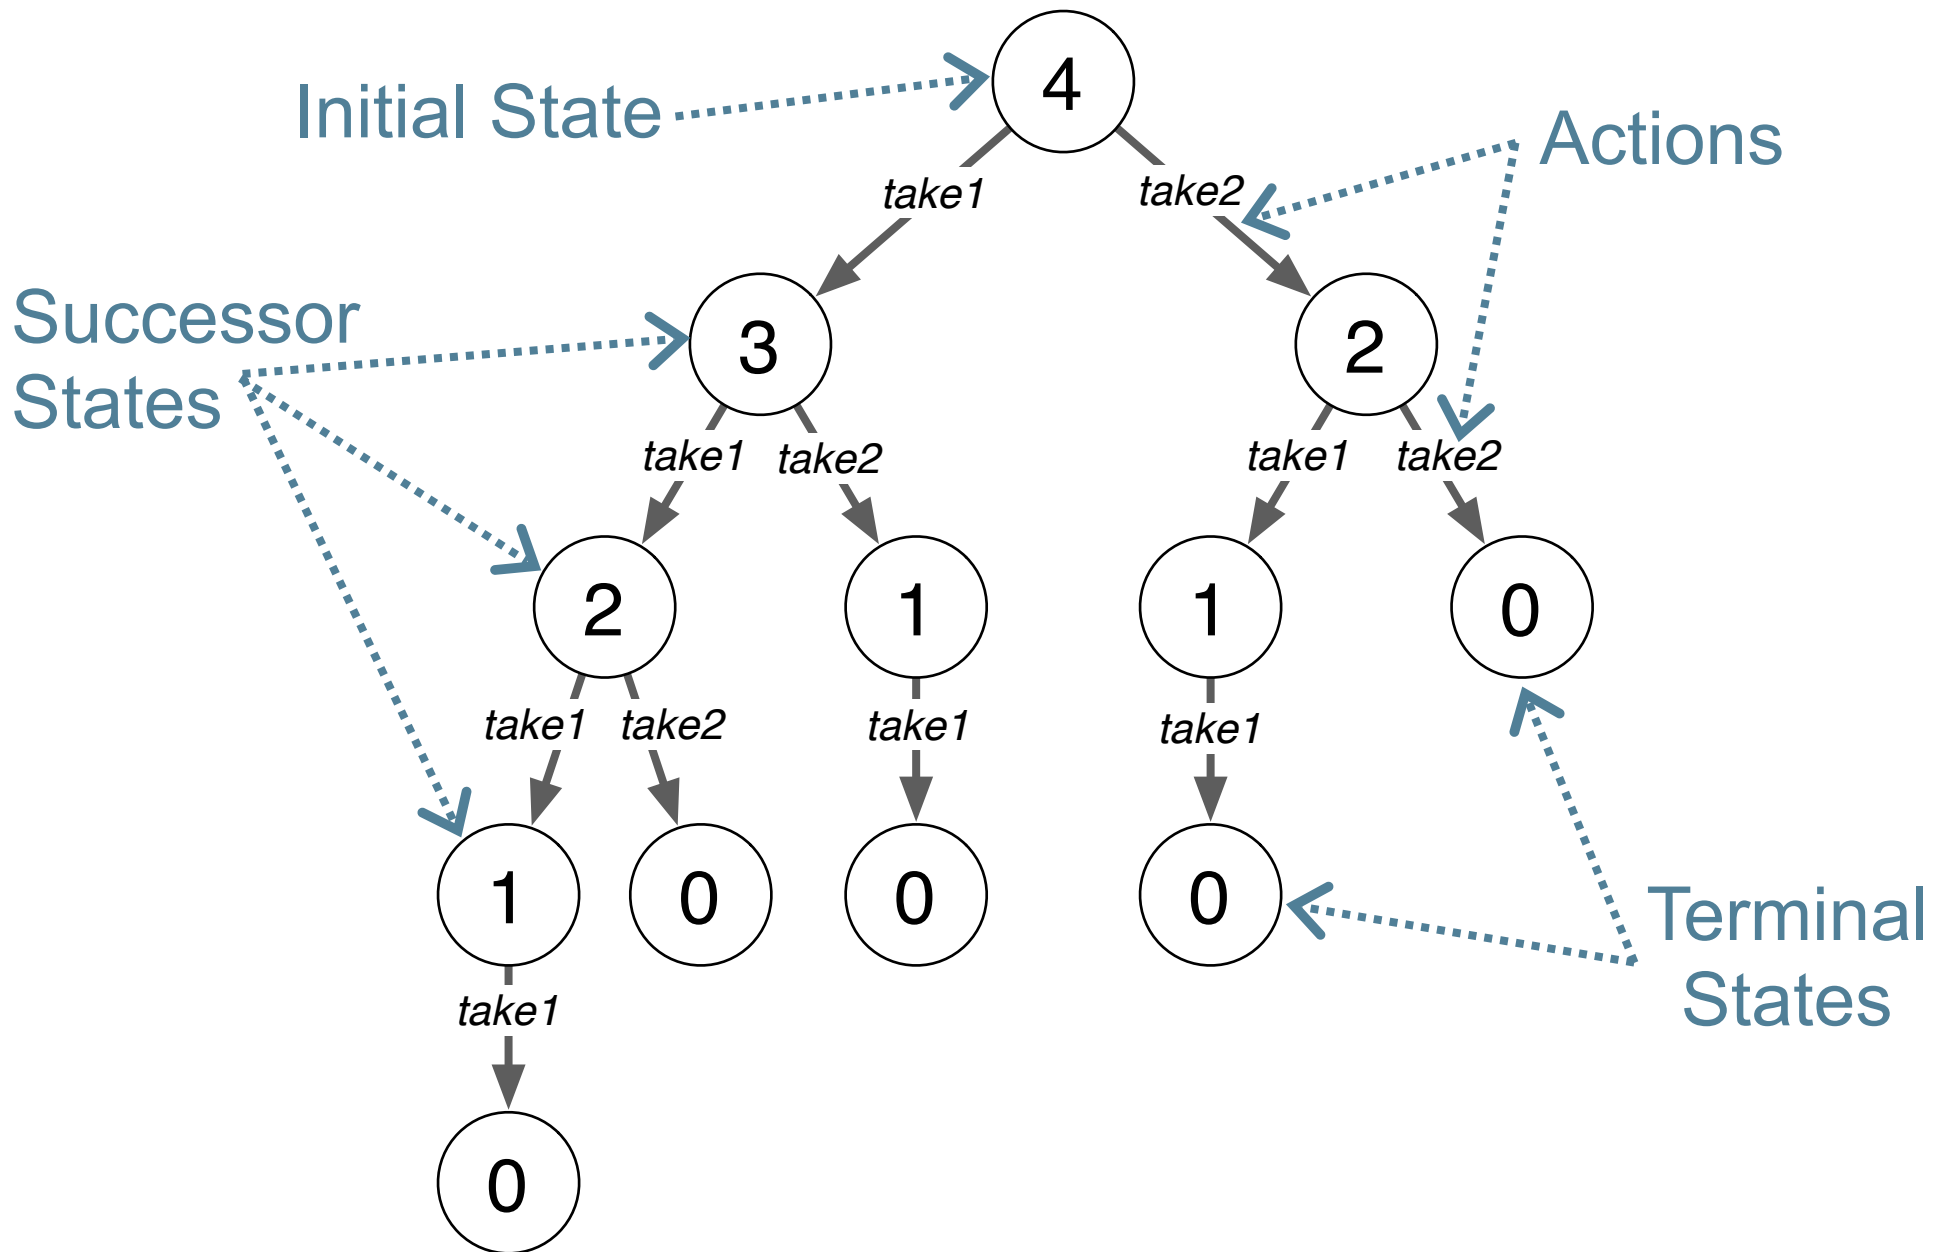
Ply a level of the search tree defined by a move by a single player

Minimax Value the value of a node in the game tree for a given player, assuming that both players play optimally to the end of the game


4-Nim State Space





4-Nim State Space



4-Nim State Space

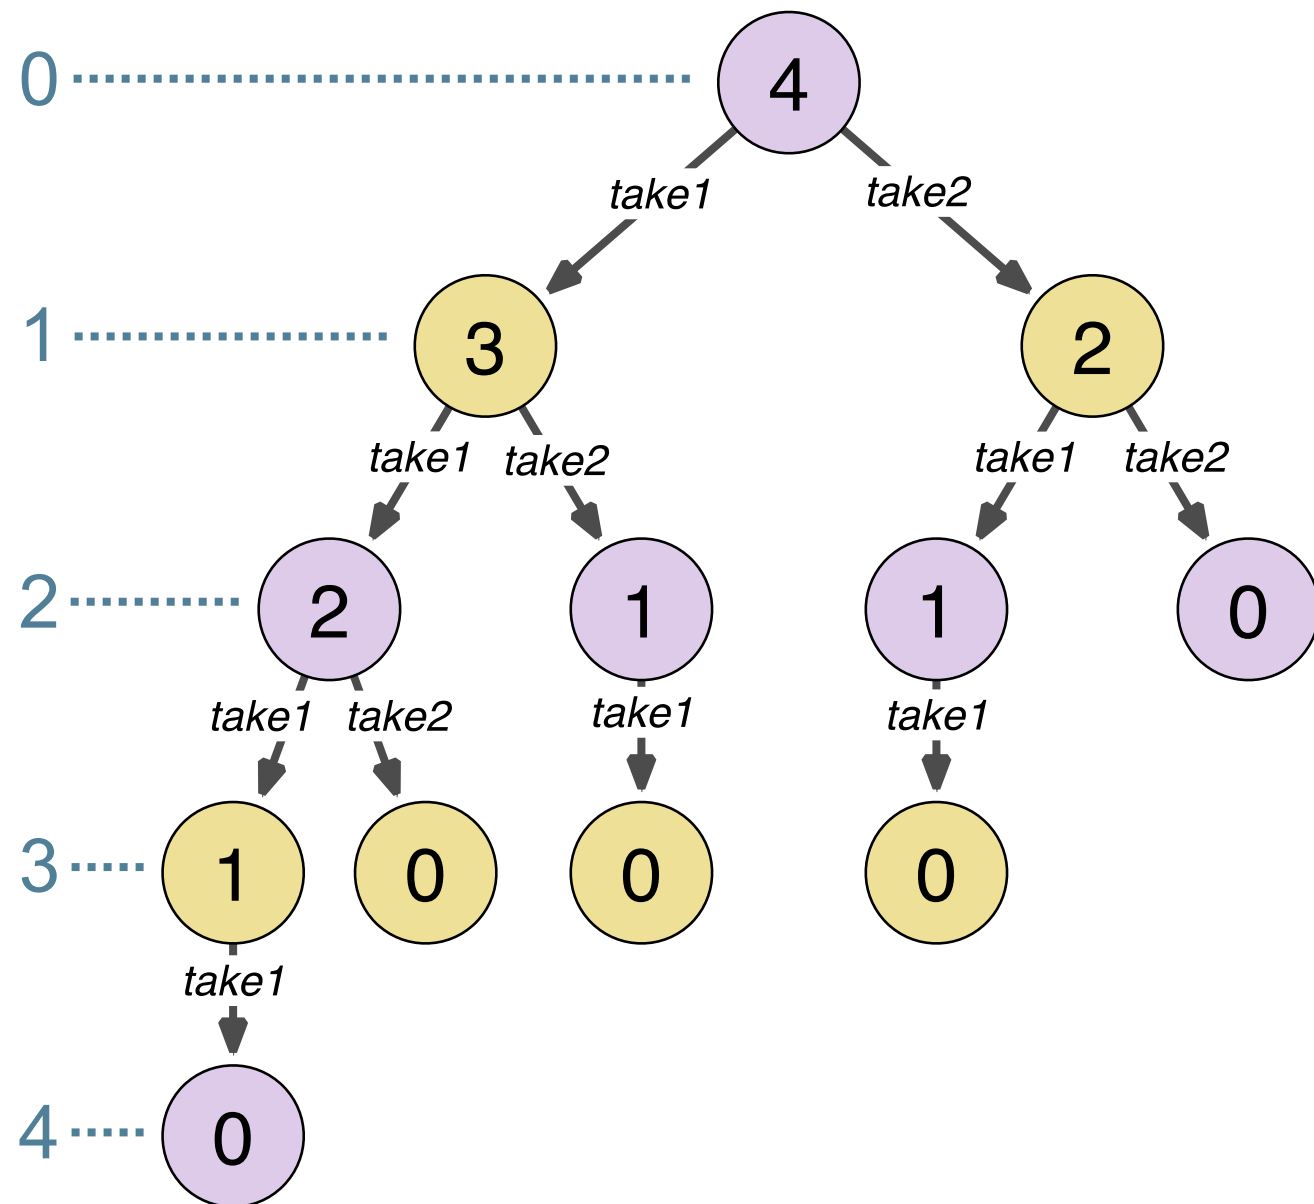


4-Nim State Space



Game Ply (4-Nim)

The **ply** (depth) indicates whose turn it is to move



Utilities (4-Nim)

The amount of reward for each terminal state is captured by the **utility function**.

P_1 Utility:

Win = 1

Loss = -1

Draw = 0

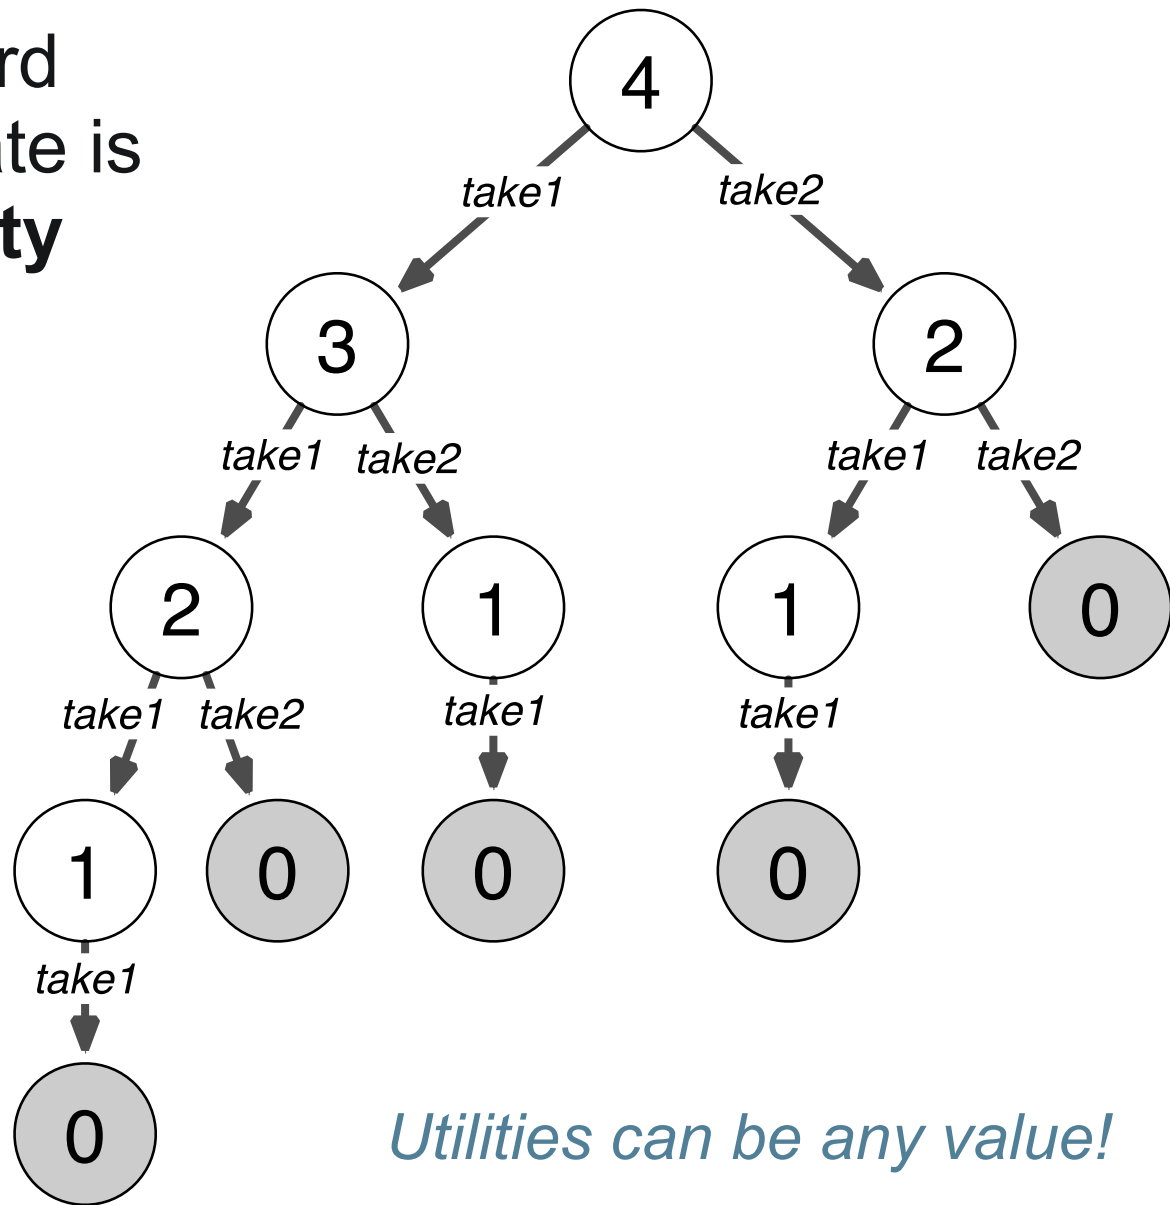
P_2 Utility:

Win = 1

Loss = -1

Draw = 0

$$\left. \begin{array}{l} U(P_1) = 1 \\ U(P_2) = -1 \end{array} \right\}$$



Utilities can be any value!

Utilities (4-Nim)

The amount of reward for each terminal state is captured by the **utility function**.

P_1 Utility:

Win = 1

Loss = -1

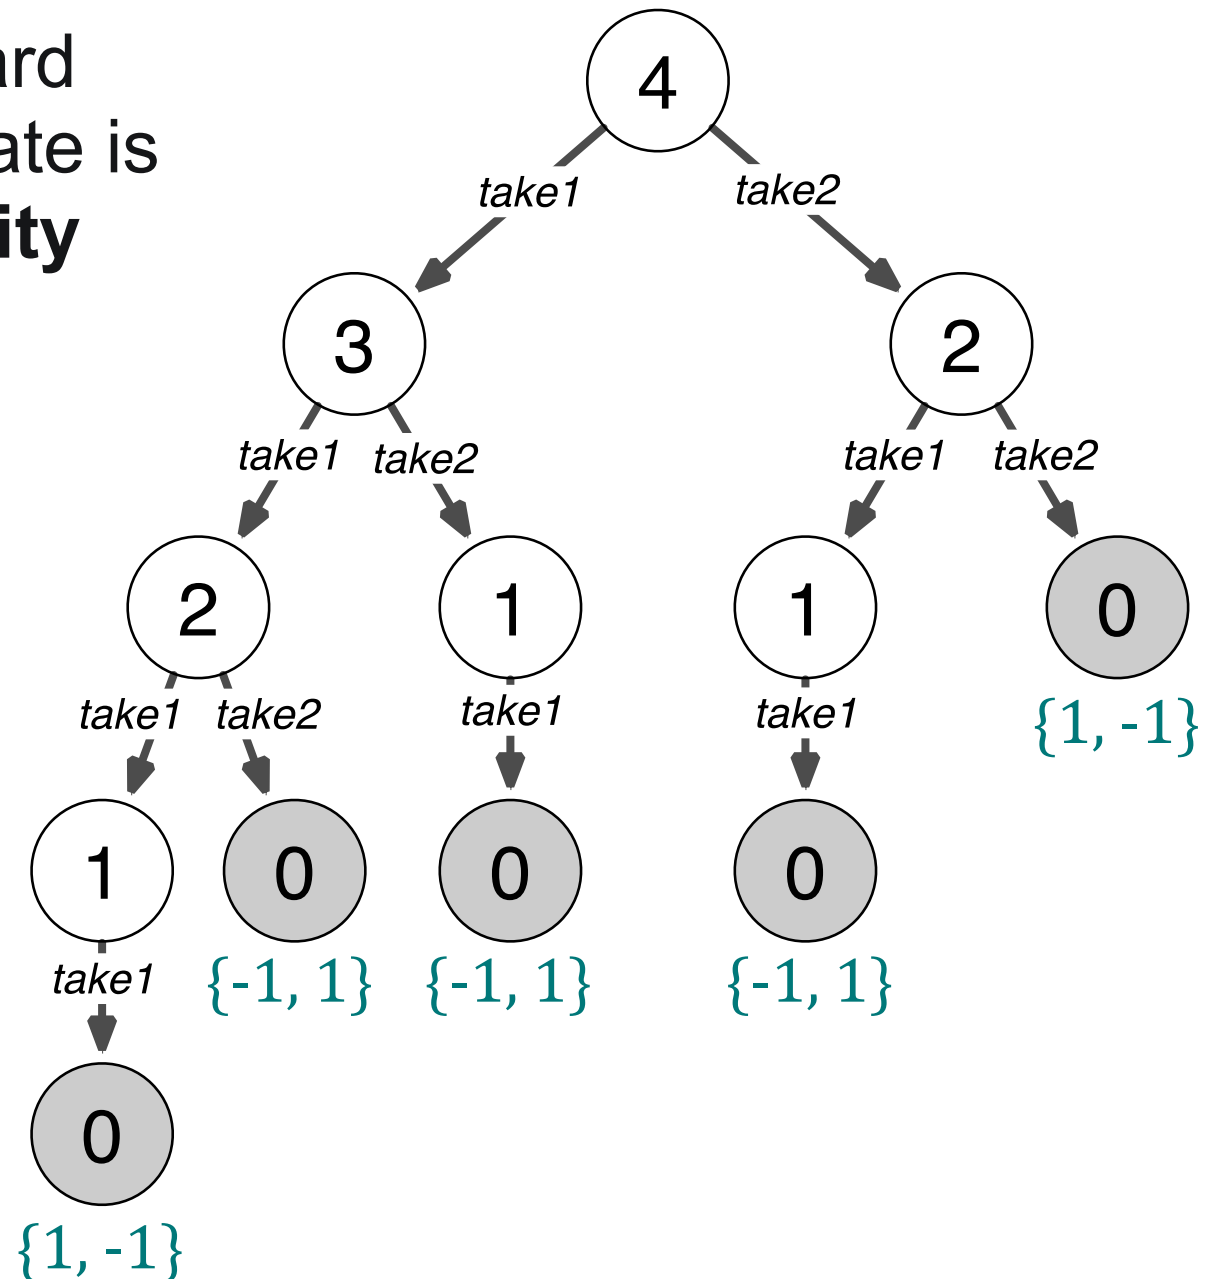
Draw = 0

P_2 Utility:

Win = 1

Loss = -1

Draw = 0

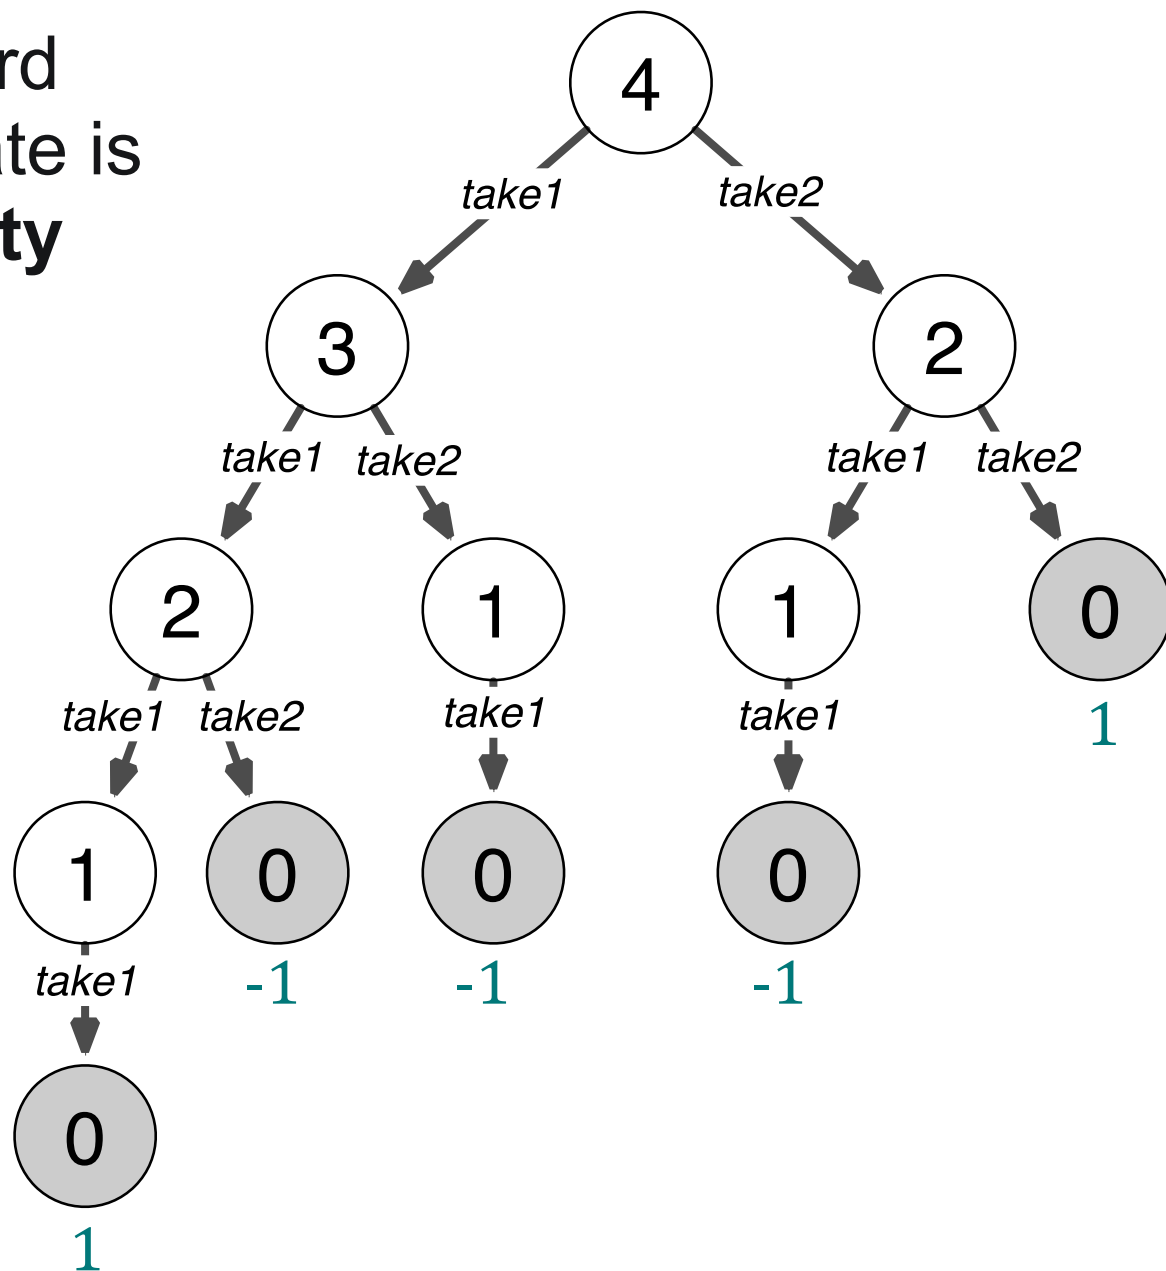


Utilities (4-Nim)

The amount of reward for each terminal state is captured by the **utility function**.

P_1 Utility:	P_2 Utility:
Win = 1	Win = 1
Loss = -1	Loss = -1
Draw = 0	Draw = 0

If game is **zero sum**, then we only need to show one utility — the one for Player 1.



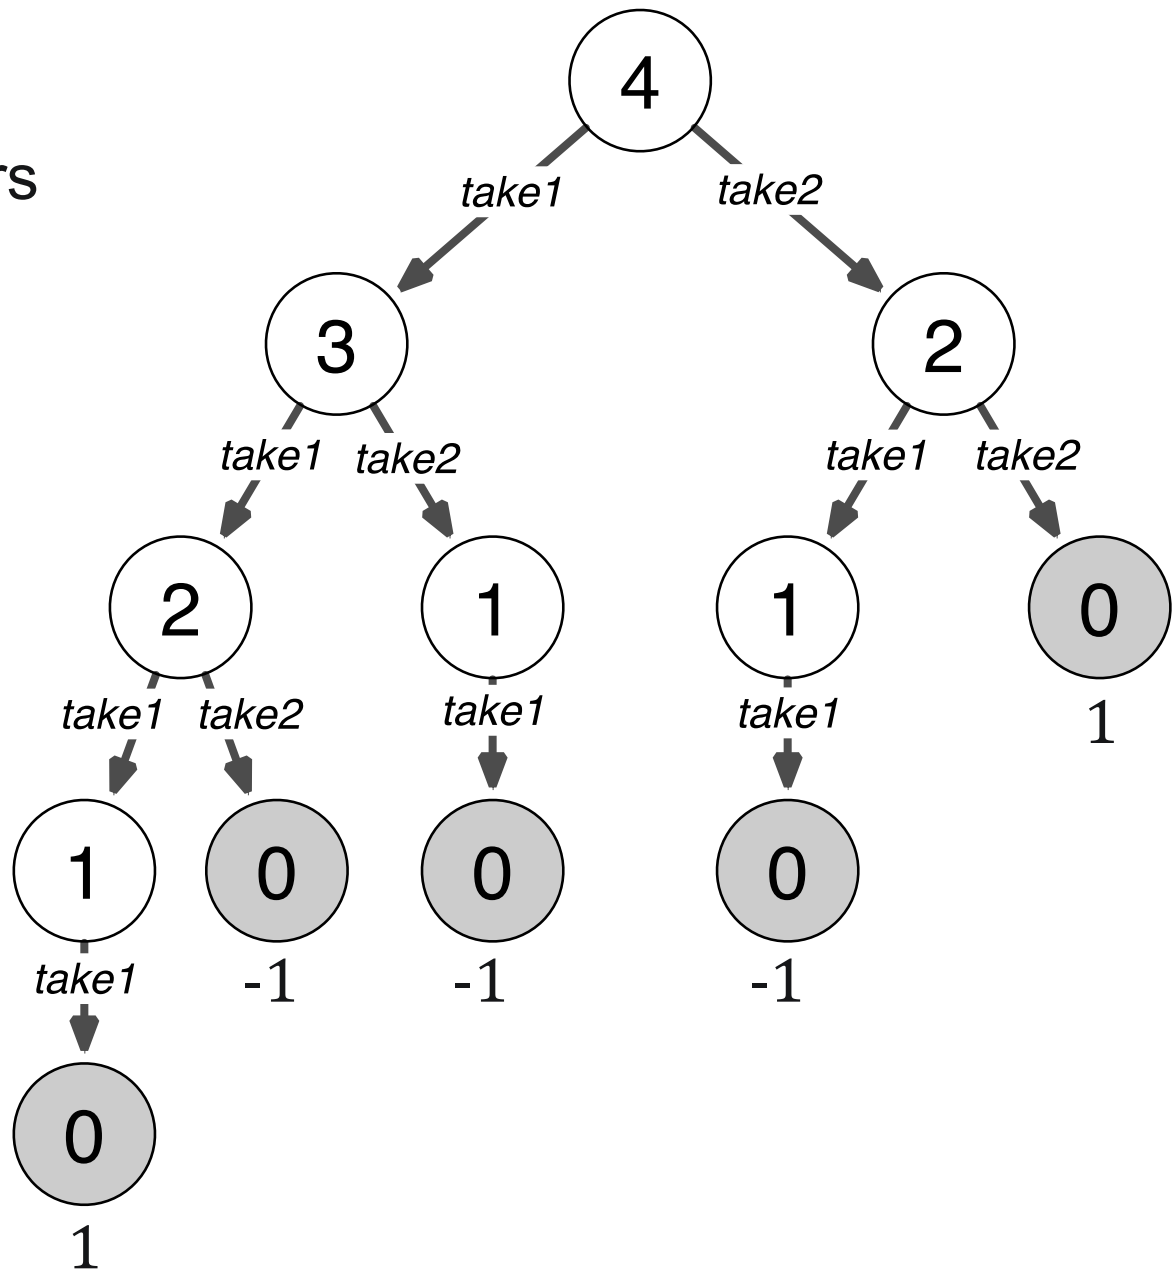
Minimax (4-Nim)

If we assume both players play optimally:

P_1 will **maximize** $U(P_1)$

P_2 will **minimize** $U(P_1)$

We can compute the expected value of an interior node with the **Minimax** algorithm.



*How can we determine the utility
of non-terminal states?*

Minimax

Minimax

- The Minimax algorithm allows us to determine the utility value of any state in the game tree.
- Minimax assumes that your opponent will behave rationally (which might not be true).
- By selecting the action leading to a state with the highest value, we can maximize our expected payoff.

Minimax Algorithm Pseudocode

function MINIMAX-DECISION(*state*) **returns** *an action*

inputs: *state*, current state in game

return the *a* in ACTIONS(*state*) maximizing MIN-VALUE(RESULT(*a*, *state*))

function MAX-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do** $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$


for *a, s* in SUCCESSORS(*state*) **do** $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

Minimax (4-Nim)

Strategy:
select moves
with highest
minimax value.

That is, select
the best
achievable
payoff against
best play by
your opponent.



Minimax (4-Nim)

Strategy:
select moves
with highest
minimax value.

MAX


MIN

That is, select
the best
achievable
payoff against
best play by
your opponent.

MAX

MIN

We use the
algorithm to
determine the
utilities of all
states.



Click time!

What does the game tree tell us about 4-Nim strategy?

(A) Player 1 has an advantage

MAX

(B) Player 2 has an advantage


MIN

(C) A priori, neither player has an advantage

MAX

(D) Most games will end in a tie

MIN



A STRANGE GAME.
THE ONLY WINNING MOVE IS
NOT TO PLAY.

HOW ABOUT A NICE GAME OF CHESS?

Properties of Minimax

Complete? Yes, if tree is finite

Optimal? Yes, against optimal opponent

Time $O(b^m)$


Space $O(bm)$ (depth-first)

b = branching factor, m = max height of tree

...but for chess, $b \approx 35$, $m \approx 100$, so *an exact solution is completely infeasible!*

Game Trees Can Be Big


- Checkers has a branching factor $b \approx 35$, with $m \approx 50$ yielding 500 billion billion game states.
- Chess has a branching factor of $b \approx 35$, with moves $m \approx 100$.
- Go has a branching factor $b \approx 250$, with $m \approx 200$
- For many of these games, Minimax traversing the entire game tree is impossible.



Improving Minimax: α - β Pruning

- α is the best value (for MAX player) found so far on this path
- β is the best value (for MIN player) found so far on this path
- If we are searching a MAX node and we find even one child with value $> \beta$, then our parent MIN node will never choose this action so we should **prune**
- If we are searching a MIN node and we find even one child with value $< \alpha$, then our parent MAX node will never choose this action so we should **prune**

α - β Pruning




α - β Pruning

MAX

$[-\infty, +\infty]$

MIN


$[-\infty, 3]$



α - β Pruning

MAX


MIN



α - β Pruning

MAX


MIN



α - β Pruning

MAX


MIN



α - β Pruning

MAX

MIN



What just happened?

- α is the value of the best (highest-value) choice found so far at any choice point along the path for MAX
- If v is worse than α , MAX will avoid it, so that branch can be pruned.
- β is the value of the best (lowest-value) choice found so far at any choice point along the path for MIN
- If v is worse than β , MIN will avoid it, so that branch can be pruned.

α - β Pruning

- Pruning produces results that are exactly equivalent to complete (unpruned) search.
- Entire subtrees can be pruned.
- Node ordering can improve effectiveness. Perfect ordering gives complexity $O(b^{m/2})$, since branching factor goes from b to \sqrt{b} . Thus α - β can search twice as far in equal time.
- You can avoid recomputing the value of previously seen states by storing them in a **transposition table**.

Minimax with Lookahead

- Even with pruning, full exploration of the game tree is often intractable (tic-tac-toe has $9! = 362,880$ states, chess has more than 10^{40} states!)
- Alternatively, you can stop the search before you reach terminal states using a cutoff test
- At some **lookahead limit** (depth limit), evaluate nodes using an **evaluation function** instead of the normal utility function
- Use the evaluation function values as if they were the true state utility values

Evaluation Functions

- Must be efficient to calculate
- Must order terminal states in the same way as the true utility function
- Typically calculate **features** – simple characteristics of the game that are strongly correlated with the probability of winning
- The evaluation function combines feature to produce a score. For example:

$$Eval(x) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s) = \sum_{i=1}^n w_i f_i(s)$$

(linear equations assume feature independence, but in practice it often doesn't matter)

Example Chess Features

- Total number of pieces
- Relative number of bishops, knights, rooks and/or pawns
- Distance of furthest pawn from start
- Relative freedom (number of possible moves)
- Other binary indicators (Has queen? Castled? In check?)

Games with Chance Elements

- For games with chance elements, we construct a game tree with a separate ply to represent the stochastic moves
- To determine the values in such a game tree, we use a variant of the Minimax algorithm called **Expectiminimax**
- Rather than minimizing or maximizing, we determine the *expected value* of chance nodes



Expected Value

The sum of the probability of each possible outcome multiplied by its value:

$$E(X) = \sum_i p_i x_i$$



$$\begin{aligned} E(\text{roll}) &= \left(\frac{1}{6} * 1\right) + \left(\frac{1}{6} * 2\right) + \left(\frac{1}{6} * 3\right) + \left(\frac{1}{6} * 4\right) + \left(\frac{1}{6} * 5\right) + \left(\frac{1}{6} * 6\right) \\ &= \frac{1}{6} + \frac{1}{3} + \frac{1}{2} + \frac{2}{3} + \frac{5}{6} + 1 = 3.5 \end{aligned}$$

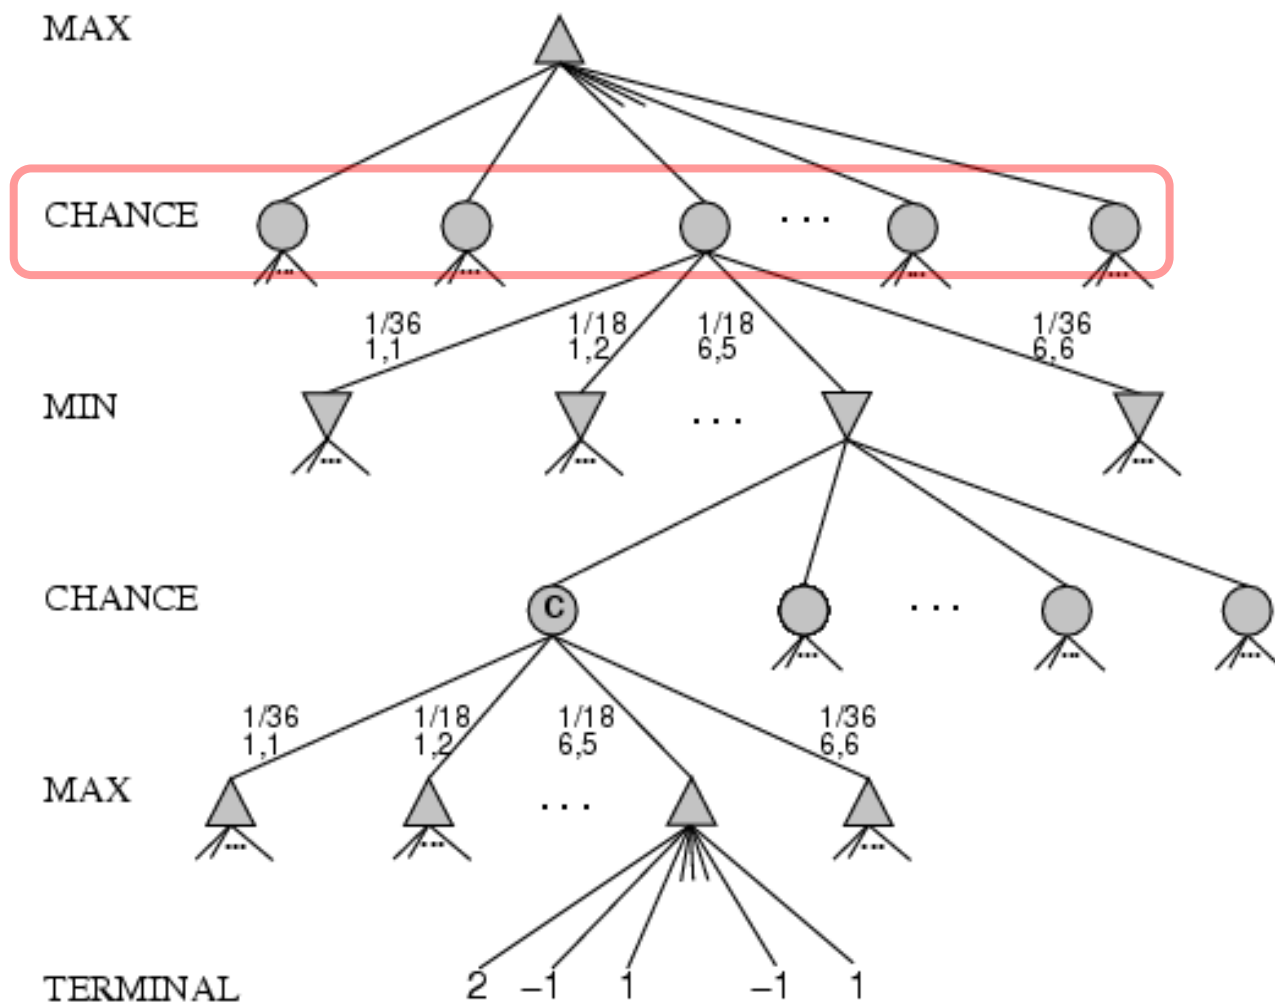
Expected Value

- The sum of the probability of each possible outcome multiplied by its value:

$$E(X) = \sum_i p_i x_i$$

- There are pathological cases where this statistic could do something strange
 - Extreme values (“outliers”)
 - Functions that are a non-linear transformation of the probability of winning

Game Tree with Chance Elements



Three different cases to evaluate, rather than just two.

Expectiminimax

EXPECTIMINIMAX(n) =

If terminal node, UTILITY(n)

If MAX node, $\max_{s \in \text{successors}(n)} \text{MINIMAX}(s)$


If MIN node, $\min_{s \in \text{successors}(n)} \text{MINIMAX}(s)$

If CHANCE node, $\sum_{s \in \text{successors}(n)} P(s) \times \text{EXPECTIMINIMAX}(s)$

Dropsy Nim

- We can add an element of stochasticity to Nim, where the results of certain actions are uncertain
- “Dropsy Nim” is the same as regular Nim, except when a player picks up two coins, they may only get one, according to some fixed probability (for example, 0.2)

Minimax (4-Nim)






Expectiminimax (Dropsy 4-Nim)

MAX

$$E(X) = \sum_i p_i x_i$$

MIN

$$E(X) = \sum_i p_i x_i$$



What's Next?

For Monday:

- Constraint Satisfaction Problems
- Read Russell & Norvig (AIMA)
Sections 6.1-6.5

THIS SLIDE INTENTIONALLY LEFT BLANK